

Boundary Information Geometry for S-box Side-Channel Leakage: Walsh Covariance, Active Fisher Structure, and CPA Belief Dynamics

Abstract

Side-channel attacks on lightweight block ciphers often reduce locally to distinguishing key hypotheses through leakage generated by bijective S-boxes. This paper develops an information-geometric model for these finite key-induced distributions while separating three security-relevant spaces: the singular distributions induced by exact key hypotheses, the regular full-support exponential family used for Fisher geometry, and the adversary’s posterior belief simplex. For an s -bit bijective S-box S and key hypothesis k , the exact distribution $P_k(a, b) = 2^{-s} \mathbf{1}[b = S(a \oplus k)]$ is supported on one graph inside $\{0, 1\}^{2s}$ and therefore lies on the boundary of the probability simplex. Walsh characters give the key-translation rule $\hat{P}_k(\alpha, \beta) = (-1)^{\alpha \cdot k} \hat{P}_0(\alpha, \beta)$. Using this rule as a structural input, we derive closed-form boundary covariance matrices and prove that, for every bijective S-box, their active subspace has dimension $2^s - 1$ and active eigenvalues equal 2^s in the unnormalised Walsh basis. We then interpret correlation power analysis under Gaussian leakage as Bayesian dynamics on the key-belief simplex. The Hamming-weight variance $s/(4\sigma^2)$ is a bijectivity-driven baseline, whereas pairwise likelihood gaps determine key-hypothesis discrimination and worst-case posterior-odds growth. Reproducible computations and Monte Carlo CPA simulations for PRESENT, GIFT, and SKINNY show how S-boxes with

001
002
003
004
005
006
007
008
009
010
011
012
013
014
015
016
017
018
019
020
021
022
023
024
025
026
027
028
029
030
031
032
033
034
035
036
037
038
039
040
041
042
043
044
045
046

the same Walsh-amplitude distribution can have distinct hard and easy key differences under Hamming-weight leakage while sharing the boundary covariance invariants forced by bijectivity.

Keywords: cybersecurity, side-channel analysis, correlation power analysis, S-box, information geometry, Fisher–Rao metric, Walsh–Hadamard transform, Bayesian inference

1 Introduction

1.1 Motivation

Side-channel attacks use implementation leakage in addition to the abstract cipher description. In many attacks on lightweight block ciphers, the local statistical object is a keyed S-box computation: an adversary observes or controls an input nibble or byte, measures a noisy leakage trace, and compares key hypotheses through predicted intermediate values. Correlation power analysis (CPA) is the standard example: each trace updates the adversary’s belief about the key using a leakage model such as Hamming weight.

The exact distribution induced by a keyed bijective S-box is not a smooth interior point of a probability simplex. For an s -bit bijective S-box $S : \{0, 1\}^s \rightarrow \{0, 1\}^s$ and key hypothesis $k \in \{0, 1\}^s$, the pair (a, b) with uniform a and $b = S(a \oplus k)$ is supported on only 2^s points of the 2^{2s} -point input-output space. Hence P_k is a singular boundary distribution. This matters because ordinary Fisher–Rao tensors, geodesics, and dual-flat coordinates are defined on full-support statistical models, not directly at support-deficient endpoints.

This paper develops a boundary information-geometric framework for this side-channel setting. The central discipline is to keep separate three objects that are often conflated in informal geometric descriptions:

- (i) the regular full-support Walsh exponential family, where ordinary Fisher geometry is valid;
- (ii) the finite boundary orbit of exact key-induced S-box distributions;
- (iii) the adversary’s posterior belief simplex over key hypotheses, where CPA updates take place.

This separation makes it possible to use information geometry in a controlled way: boundary distributions are represented through covariance limits, whereas Bayesian key recovery remains a regular trajectory on the belief simplex as long as the prior has full support and the leakage likelihood is positive.

For cybersecurity and hardware-security evaluation, the practical output of this separation is a diagnostic workflow for identifying which quantities are unavoidable consequences of bijectivity, which quantities are artefacts of the chosen Walsh coordinates, and which quantities actually affect CPA discrimination. In particular, for Hamming-weight leakage the scalar variance is identical for all bijective S-boxes of the same width, but the pairwise gaps between key hypotheses need not be identical. Those gaps determine the expected posterior-odds growth against wrong keys and therefore identify the key differences that are locally hardest or easiest for a CPA adversary to reject.

1.2 Relation to prior work

The information-geometric background is the Fisher–Rao and dual-affine geometry of finite statistical models and regular exponential families [1–3]. The closest coding-theoretic precedent is the work of Ikeda, Tanaka, and Amari on LDPC and turbo codes [4]: parity-check statistics define a regular information-geometric model and




Fig. 1 Boundary-to-belief framework. The exact keyed S-box laws lie on a singular boundary orbit, the Fisher calculations are made in a regular Walsh exponential family or in boundary covariance limits, and CPA updates evolve on the separate key-belief simplex.

decoding trajectories are interpreted through projections. Here, the sufficient statistics are Walsh–Hadamard characters on S-box input-output pairs, and the inference trajectory is the CPA posterior over key hypotheses.

On the cryptographic side, the work is related to classical Walsh/LAT and DDT analysis of vectorial Boolean functions and S-boxes, including Nyberg’s differential uniformity [5]. Four-bit S-boxes and optimal 4-bit permutations have been classified in the Boolean-function literature, notably by Leander and Poschmann [6]. Recent geometric viewpoints in symmetric cryptanalysis also include geometric formulations of linear and differential cryptanalysis [7, 8]. The present paper builds on this background by asking which information-geometric objects are associated with exact keyed S-box distributions and how those objects connect to side-channel belief updates. This positioning uses existing 4-bit S-box classifications as prior cryptographic knowledge and focuses on the local leakage distributions used in CPA.

On the side-channel side, the leakage model and posterior update are consistent with differential and correlation power analysis [9, 10] and with information-theoretic analyses of key recovery attacks [11, 12]. The present framework gives a formal model

Table 1 Relation to existing metrics.

Metric or framework	Standard role	Role in this paper
Walsh/LAT spectrum	Measures linear correlations of vectorial Boolean functions and supports nonlinearity analysis.	Provides coordinates for the key-induced boundary orbit and for closed-form covariance diagnostics.
DDT and differential uniformity	Measures differential propagation and classical S-box resistance to differential attacks [5].	Used as part of the surrounding S-box context for separating classical cryptographic criteria from leakage-geometry diagnostics.
4-bit S-box classification	Classifies small permutations up to established equivalences [6].	Used as prior cryptographic knowledge for interpreting the local leakage geometry.
CPA and DPA likelihood models	Explain key recovery from noisy implementation leakage [9, 10].	Interpreted as Bayesian dynamics on the key-belief simplex with pairwise likelihood gaps.
Mutual information and key-rank measures	Quantify leakage and attack success at the experiment or trace-distribution level [11, 12].	Complemented by local structural diagnostics for keyed S-box laws and posterior-odds growth.
Information geometry of codes	Studies decoding and projections in regular statistical models [4].	Adapted to singular S-box graph distributions by separating boundary covariance, regular interiors, and belief-simplex updates.

for separating baseline invariants forced by bijectivity from pairwise leakage gaps that control key-hypothesis discrimination.

Table 1 summarizes the intended position of the framework relative to standard S-box and side-channel criteria.

1.3 Setting and notation

Throughout, $s \geq 1$ denotes the S-box width in bits, $S : \{0, 1\}^s \rightarrow \{0, 1\}^s$ is a fixed bijection, $k \in \{0, 1\}^s$ is a key hypothesis, and $a \in \{0, 1\}^s$ is a uniformly distributed plaintext nibble or byte. All additions inside bit strings are over $\mathbb{F}_2$. The inner product is $\alpha \cdot a = \bigoplus_{i=1}^s \alpha_i a_i$. For $(\alpha, \beta) \in \{0, 1\}^s \times \{0, 1\}^s$, the Walsh character is

$$\chi_{\alpha\beta}(a, b) = (-1)^{\alpha \cdot a \oplus \beta \cdot b}.$$

We write $N = 2^s$ and $\mathcal{X} = \{0, 1\}^s \times \{0, 1\}^s$.

231 1.4 Contributions

232

233 **C1. Boundary model for key-induced S-box distributions.** We model the exact
 234
 235 keyed S-box law

$$236 P_k(a, b) = 2^{-s} \mathbf{1}[b = S(a \oplus k)]$$

237
 238 as a boundary distribution on the input-output simplex. The Walsh translation
 239
 240 rule

$$241 \hat{P}_k(\alpha, \beta) = (-1)^{\alpha \cdot k} \hat{P}_0(\alpha, \beta)$$

242
 243 is used as a coordinate identity for the finite key orbit and as the algebraic input
 244
 245 for the covariance and leakage diagnostics.
 246

247 **C2. Closed-form boundary covariance and active Fisher structure.** We derive
 248
 249 the boundary covariance form

250

251

$$252 \mathcal{I}_{ij}^\partial(k) = \text{Cov}_{P_k}(T_i, T_j)$$

253

254

255 in Walsh coordinates and prove that, for every bijection $S : \{0, 1\}^s \rightarrow \{0, 1\}^s$, this
 256
 257 matrix has rank $2^s - 1$ and active eigenvalues all equal to 2^s in the unnormalised
 258
 259 Walsh basis. Thus the active boundary geometry is a structural property of
 260
 261 bijective S-box graphs shared by the three illustrative 4-bit examples.

262 **C3. Regular interior and belief-simplex geometry.** We explicitly separate the
 263
 264 regular Walsh exponential family, where the usual Amari dual-flat structure and
 265
 266 Pythagorean identities apply, from the singular boundary orbit. The adversary's
 267
 268 key belief is treated as a separate full-support simplex, equipped with its own
 269
 270 Fisher–Rao metric.

271 **C4. CPA as Bayesian belief dynamics with pairwise leakage gaps.** Under
 272
 273 Gaussian leakage, each CPA trace is represented as a Bayesian update on the key-
 274
 275 belief simplex. The Hamming-weight variance $s/(4\sigma^2)$ is identified as a baseline
 276
 forced by bijectivity and uniform inputs. The quantities that distinguish key

hypotheses are the pairwise separations	277
	278
	279
$\Delta(k, k^*) = \frac{1}{2^s} \sum_a (f(S(a \oplus k^*)) - f(S(a \oplus k)))^2,$	280
	281
	282
which determine expected log-Bayes gains.	283
	284
C5. Reproducible diagnostics on lightweight S-boxes. For PRESENT, GIFT,	285
and SKINNY, we compute Walsh spectra, boundary covariance matrices, local	286
spectral separation proxies, pairwise Hamming-weight leakage gaps, Monte Carlo	287
CPA success curves, and the projected active-eigenspace curvature diagnostic.	288
	289
The examples are used to separate invariants forced by bijectivity or by the shared	290
Walsh amplitude distribution from quantities that can vary across S-boxes. The	291
pairwise-gap tables and CPA simulations identify the hard key differences and	292
trace-dependent key-rank behavior under the assumed leakage model.	293
	294
	295
	296
	297
	298
	299
1.5 Model coverage	300
	301
The paper focuses on isolated bijective S-boxes under uniform inputs and on an additive	302
Gaussian leakage model for CPA belief updates. This coverage matches the local	303
leakage setting in which S-box output predictions drive non-profiled CPA. Extensions	304
to complete multi-round ciphers, masking, higher-order leakage, non-Gaussian leakage,	305
profiled attacks, and implementation-specific trace effects can be built by adding the	306
corresponding leakage and cipher propagation models to the boundary-distribution	307
and belief-update framework.	308
	309
	310
	311
	312
	313
	314
	315
	316
	317
	318
	319
	320
	321
	322

323 2 Background

324

325

326

2.1 S-boxes and the Walsh–Hadamard transform

327

328 An s -bit S -box is a bijection $S : \{0, 1\}^s \rightarrow \{0, 1\}^s$. The i -th output bit is a Boolean

329

330

function $S_i : \{0, 1\}^s \rightarrow \{0, 1\}$.

331

332

333

334

335

336

Definition 2.1 (Walsh–Hadamard Transform) The Walsh–Hadamard transform (WHT) of

337

S at $(\alpha, \beta) \in \{0, 1\}^s \times \{0, 1\}^s$ is:

338

339

340

$$\widehat{S}(\alpha, \beta) = \sum_{a \in \{0, 1\}^s} (-1)^{\alpha \cdot a \oplus \beta \cdot S(a)} \in \{-2^s, \dots, +2^s\}, \quad 2 \mid \widehat{S}(\alpha, \beta). \quad (1)$$

341

342

By convention $\widehat{S}(0, 0) = 2^s$.

343

344

345

346

347

348

The WHT encodes the two fundamental cryptographic properties of S :

349

350

351

352

• **Nonlinearity:** $\text{NL}(S) = 2^{s-1} - \frac{1}{2} \max_{(\alpha, \beta) \neq (0, 0)} |\widehat{S}(\alpha, \beta)|$.

353

354

• **Differential uniformity:** $\delta(S) = \max_{\Delta_{\text{in}} \neq 0, \Delta_{\text{out}}} |\{a : S(a) \oplus S(a \oplus \Delta_{\text{in}}) = \Delta_{\text{out}}\}|$.

355

356

357

The *Linear Approximation Table* (LAT) entry is $\text{LAT}_S(\alpha, \beta) = \widehat{S}(\alpha, \beta)/2$. For any

358

S -box, Parseval's identity gives:

359

360

361

362

363

$$\sum_{(\alpha, \beta) \in \{0, 1\}^s \times \{0, 1\}^s} \widehat{S}(\alpha, \beta)^2 = 2^{2s} \cdot 2^s = 2^{3s}. \quad (2)$$

364

365

Example 2.1 (PRESENT S-box) The PRESENT-80 S-box ($s = 4$) is the hexadecimal table

366

367

$S = [\text{C}, 5, 6, \text{B}, 9, 0, \text{A}, \text{D}, 3, \text{E}, \text{F}, 8, 4, 7, 1, 2]$. Its nonlinearity is $\text{NL}(S) = 4$ and differential

368

uniformity $\delta(S) = 4$.

2.1.1 Remark on 4-bit S-box classification

The 4-bit case is small enough that strong classification results are available. In particular, optimal 4-bit S-boxes have been studied up to the standard equivalence relations used for vectorial Boolean functions [6]. The numerical examples in Section 7 use this classification background to provide reproducible side-channel diagnostics for widely used lightweight S-boxes and illustrate which information-geometric quantities are forced by bijectivity or by common Walsh spectra.

2.2 Information geometry: essentials

We collect the minimal background from Amari and Nagaoka [1] and Ay et al. [2] required in the sequel.

2.2.1 Statistical manifolds and the Fisher metric

Let $\mathcal{X}$ be a finite set with $|\mathcal{X}| = N$. The *probability simplex* is $\Delta_{N-1} = \{p \in \mathbb{R}^N : p_x \geq 0, \sum_x p_x = 1\}$. Its interior $\mathring{\Delta}_{N-1}$ is an $(N-1)$ -dimensional manifold.

Definition 2.2 (Fisher–Rao Metric) The *Fisher–Rao metric* on $\mathring{\Delta}_{N-1}$ at p is the positive definite bilinear form

$$g_p(u, v) = \sum_{x \in \mathcal{X}} \frac{u_x v_x}{p_x}, \quad u, v \in T_p \mathring{\Delta}_{N-1} = \left\{ w \in \mathbb{R}^N : \sum_x w_x = 0 \right\}. \quad (3)$$

By Chentsov’s theorem [3], g_F is the unique Riemannian metric on $\mathring{\Delta}_{N-1}$ (up to a positive constant) that is invariant under sufficient statistics, making it the canonical statistical distance.

415 2.2.2 Exponential families as flat submanifolds

416

417 **Definition 2.3** (Exponential Family) Given functions $T_1, \dots, T_d : \mathcal{X} \rightarrow \mathbb{R}$ (the *sufficient*
418 *statistics*), the *exponential family* $\mathcal{E}$ is:

419

420

421

422

$$p(x; \theta) = \exp\left(\sum_{i=1}^d \theta_i T_i(x) - \psi(\theta)\right), \quad \theta \in \Theta \subseteq \mathbb{R}^d, \quad (4)$$

423

424

425

426

427

428

429

430

431

432

433

434

435

436

437

438

439

440

441

442

443

444

445

446

447

448

449

450

451

452

453

454

455

456

457

458

459

460

The natural parameters θ_i and the expectation parameters $\eta_i = \mathbb{E}_\theta[T_i(X)]$ are related by $\eta = \nabla \psi(\theta)$ (the bijection is guaranteed when the family is minimal and regular).

Theorem 2.1 (Dually Flat Structure, [1]) *Every regular minimal exponential family $\mathcal{E}$ is a dually flat submanifold of $\mathring{\Delta}_{N-1}$: the e -connection (exponential connection) is flat in natural parameters θ , and the m -connection (mixture connection) is flat in expectation parameters η . The generalised Pythagorean theorem holds: if $P_A, P_B, P_C \in \mathcal{E}$ and the e -geodesic from A to B is g_F -orthogonal to the m -geodesic from B to C , then*

$$D_{\text{KL}}(P_A \| P_C) = D_{\text{KL}}(P_A \| P_B) + D_{\text{KL}}(P_B \| P_C). \quad (5)$$

446 2.2.3 Mixture projection and Bayesian inference

447

448

449

450

451

452

453

454

455

456

457

458

459

460

Definition 2.4 (m -Projection) Given $p \in \mathring{\Delta}_{N-1}$ and a closed m -flat submanifold $\mathcal{M} \subseteq \mathring{\Delta}_{N-1}$, the m -projection of p onto $\mathcal{M}$ is:

$$\Pi_{\mathcal{M}}^{(m)}(p) = \arg \min_{q \in \mathcal{M}} D_{\text{KL}}(q \| p). \quad (6)$$

The following classical result (see [1], Chapter 3) is central to Section 5:

Proposition 2.1 (Bayesian Update as m -Projection) *Let $p(x, y) = p(x)p(y|x)$ be a joint distribution on $\mathcal{X} \times \mathcal{Y}$. After observing $y = y_0$, the posterior $p(x|y_0)$ is the m -projection of*

the joint p onto the submanifold $\mathcal{M}_{y_0} = \{q \in \Delta : q(y) = 0 \text{ for } y \neq y_0\}$, restricted to the x -marginal.

2.2.4 The LDPC analogy: Ikeda-Tanaka-Amari (2004)

Ikeda, Tanaka, and Amari [4] develop a complete information-geometric framework for LDPC and turbo codes. They model the joint distribution over a binary codeword $x \in \{-1, +1\}^N$ as:

$$p(x; \theta, v) = \exp\left(\theta \cdot x + \sum_r v_r c_r(x) - \psi\right), \quad (7)$$

where $c_r(x) = \prod_{i \in r} x_i$ are parity check functions and v_r the constraint weights. They show: (i) the e/m dual structure is exact; (ii) Belief Propagation decoding corresponds to iterated e/m -projections; (iii) curvature of the statistical manifold controls decoding convergence.

In the present paper, the parity functions $c_r(x)$ are replaced by the Walsh-Hadamard characters $\chi_{\alpha\beta}(a, b)$ of the S-box, and LDPC decoding is replaced by CPA key recovery. Table 2 summarises the correspondence.

Table 2 Structural analogy between the LDPC framework of Ikeda-Tanaka-Amari (2004) and the S-box framework of this paper.

Ikeda-Tanaka-Amari (LDPC)	This paper (S-box)
Codeword $x \in \{-1, +1\}^N$	Input-output pair $(a, b) \in \{0, 1\}^{2s}$
Parity check $c_r(x) = \prod_{i \in r} x_i$	Walsh character $\chi_{\alpha\beta}(a, b) = (-1)^{\alpha \cdot a \oplus \beta \cdot b}$
Channel log-likelihood θ_i	Key log-likelihood ratio $\log(q_k/q_{k'})$
Parity constraint weight v_r	Walsh spectrum value $\hat{S}(\alpha, \beta)$
Belief Propagation decoding	CPA key recovery
Decoding convergence rate	Per-trace signal information $I_S(\sigma)$

507 3 The S-box Statistical Model

508

509

510

511

512

513

514

515

516

517

518

519

520

521

522

523

524

525

526

527

528

529

530

531

532

533

534

535

536

537

538

539

540

541

542

543

544

545

546

547

548

549

550

551

552

3.1 Key-induced boundary distributions

Fix a bijective S-box $S : \{0, 1\}^s \rightarrow \{0, 1\}^s$ and let A be uniformly distributed over $\{0, 1\}^s$. Under key hypothesis $k \in \{0, 1\}^s$, the observed input-output pair is $(A, S(A \oplus k))$. This induces the following distribution on $\mathcal{X} = \{0, 1\}^s \times \{0, 1\}^s$.

Definition 3.1 (Key-induced distribution) For $k \in \{0, 1\}^s$, define

$$P_k(a, b) = 2^{-s} \mathbf{1}[b = S(a \oplus k)]. \quad (8)$$

The support of P_k is the graph $G_k = \{(a, S(a \oplus k)) : a \in \{0, 1\}^s\}$.

Each P_k is a boundary point of the simplex $\Delta_{2^{2s}-1}$, because it assigns zero mass to $2^{2s} - 2^s$ outcomes. This support deficiency is central: ordinary Fisher–Rao tensors of the full simplex are defined on the interior, not directly at P_k . We therefore treat P_k as a singular endpoint of an interior regularization when differential-geometric operations are needed.

3.2 Walsh representation and key translation

Definition 3.2 (Walsh coefficient of P_k) For $(\alpha, \beta) \in \{0, 1\}^s \times \{0, 1\}^s$, set

$$\hat{P}_k(\alpha, \beta) = \sum_{(a, b) \in \mathcal{X}} \chi_{\alpha\beta}(a, b) P_k(a, b) = 2^{-s} \sum_{a \in \{0, 1\}^s} (-1)^{\alpha \cdot a \oplus \beta \cdot S(a \oplus k)}. \quad (9)$$

Lemma 3.1 (Walsh key-translation rule) For every $k \in \{0, 1\}^s$ and every $(\alpha, \beta) \in \{0, 1\}^s \times \{0, 1\}^s$,

$$\hat{P}_k(\alpha, \beta) = (-1)^{\alpha \cdot k} \hat{P}_0(\alpha, \beta) = 2^{-s} (-1)^{\alpha \cdot k} \hat{S}(\alpha, \beta). \quad (10)$$

Proof Using $a' = a \oplus k$,

$$\hat{P}_k(\alpha, \beta) = 2^{-s} \sum_{a'} (-1)^{\alpha \cdot (a' \oplus k) \oplus \beta \cdot S(a')} = 2^{-s} (-1)^{\alpha \cdot k} \sum_{a'} (-1)^{\alpha \cdot a' \oplus \beta \cdot S(a')},$$

which is exactly the claimed identity. $\square$

Remark 3.1 (Walsh orbit) The key does not change the amplitude $|\hat{S}(\alpha, \beta)|$; it only changes the sign according to the character $\alpha \mapsto (-1)^{\alpha \cdot k}$. Thus $\{P_k : k \in \{0, 1\}^s\}$ is best described as a finite Walsh-phase orbit, not as a smooth manifold. The rule is a structural translation identity; the information-geometric content in the sequel comes from the covariance limit, active subspace, and belief-update geometry built on top of this identity.

3.3 The regular Walsh exponential family

Definition 3.3 (Walsh exponential family) Let $T_{\alpha\beta} = \chi_{\alpha\beta}$ for $(\alpha, \beta) \neq (0, 0)$. The Walsh exponential family on $\mathcal{X}$ is

$$p_\theta(a, b) = \exp \left(\sum_{(\alpha, \beta) \neq (0, 0)} \theta_{\alpha\beta} T_{\alpha\beta}(a, b) - \psi(\theta) \right), \quad \theta \in \mathbb{R}^{2^{2s}-1}. \quad (11)$$

Because the nonconstant Walsh characters together with the constant character form an orthogonal basis of all real functions on $\mathcal{X}$, this is simply the full regular exponential family over $\mathcal{X}$ written in Walsh coordinates. Its image is the interior of the probability simplex. The singular distributions P_k belong to the closure, not to the regular parameter space.

Proposition 3.1 (Boundary expectation coordinates) For each k and $(\alpha, \beta) \neq (0, 0)$,

$$\eta_{\alpha\beta}^{(k)} := \mathbb{E}_{P_k}[T_{\alpha\beta}] = (-1)^{\alpha \cdot k} 2^{-s} \hat{S}(\alpha, \beta). \quad (12)$$

The corresponding natural parameters are not finite; they are obtained only as limits of full-support approximations approaching P_k .

599 *Proof* The first equality is the definition of expectation coordinates and the second is
600 Lemma 3.1. Since P_k has zeros on $\mathcal{X}$, no finite vector θ in (11) can represent it. $\square$
601

602 3.4 Interior regularization and boundary covariance form

603
604 Let U denote the uniform distribution on $\mathcal{X}$. For $\varepsilon \in (0, 1)$ define
605

$$606 P_k^\varepsilon = (1 - \varepsilon)P_k + \varepsilon U. \quad (13)$$

607
608 Then P_k^ε has full support and belongs to the regular simplex interior. Fisher–Rao
609 calculations are unambiguous at P_k^ε . The finite covariance limit as $\varepsilon \downarrow 0$ is the following.
610

611 **Definition 3.4** (Boundary Fisher covariance form) The boundary covariance form associated
612 with key k is

$$613 \mathcal{I}_{ij}^\partial(k) = \text{Cov}_{P_k}(T_i, T_j), \quad i, j \in \{0, 1\}^s \times \{0, 1\}^s \setminus \{(0, 0)\}. \quad (14)$$

614 It is positive semidefinite and generally singular. We do not call it a non-degenerate Fisher
615 metric at P_k .
616

617 **Proposition 3.2** (Closed form of the boundary covariance) For $i = (\alpha, \beta)$ and $j = (\alpha', \beta')$,
618

$$619 \begin{aligned} 620 \mathcal{I}_{ij}^\partial(k) &= \text{Cov}_{P_k}(T_{\alpha\beta}, T_{\alpha'\beta'}) \\ 621 &= (-1)^{(\alpha \oplus \alpha') \cdot k} \left(2^{-s} \widehat{S}(\alpha \oplus \alpha', \beta \oplus \beta') - 2^{-2s} \widehat{S}(\alpha, \beta) \widehat{S}(\alpha', \beta') \right). \end{aligned} \quad (15)$$

622 In particular,

$$623 \mathcal{I}_{(\alpha\beta), (\alpha\beta)}^\partial(k) = 1 - \left(2^{-s} \widehat{S}(\alpha, \beta) \right)^2. \quad (16)$$

624
625 *Proof* Since $T_i T_j = T_{i \oplus j}$,
626

$$627 \mathbb{E}_{P_k}[T_i T_j] = \widehat{P}_k(\alpha \oplus \alpha', \beta \oplus \beta') = (-1)^{(\alpha \oplus \alpha') \cdot k} 2^{-s} \widehat{S}(\alpha \oplus \alpha', \beta \oplus \beta').$$

Subtracting $\mathbb{E}_{P_k}[T_i]\mathbb{E}_{P_k}[T_j]$ and applying Lemma 3.1 gives (15). The diagonal follows from $T_i^2 = 1$. $\square$

Lemma 3.2 (Key covariance isometry) *Let D_k be the diagonal sign matrix $(D_k)_{(\alpha\beta),(\alpha\beta)} = (-1)^{\alpha \cdot k}$. Then*

$$\mathcal{I}^\partial(k) = D_k \mathcal{I}^\partial(0) D_k. \quad (17)$$

Thus every spectral invariant of the boundary covariance form is independent of the true key.

Proof Equation (15) factors exactly into the sign contribution $(-1)^{\alpha \cdot k} (-1)^{\alpha' \cdot k}$ times the $k = 0$ covariance entry. $\square$

Remark 3.2 (What is metric and what is not) For $\varepsilon > 0$, the Fisher–Rao metric at P_k^ε is an ordinary Riemannian metric. At $\varepsilon = 0$, only the finite covariance form (14) is used. Distances or curvatures involving the exact P_k are therefore either (i) computed on the belief simplex, (ii) computed for P_k^ε and then studied as limits, or (iii) explicitly described as first-order/local covariance approximations.

3.5 First-order spectral separation of key hypotheses

The exact distributions P_k and $P_{k'}$ have disjoint supports when $k \neq k'$, so their KL divergence is infinite. Nevertheless, the Walsh coordinate displacement has a simple closed form:

$$\eta_{\alpha\beta}^{(k')} - \eta_{\alpha\beta}^{(k)} = \left((-1)^{\alpha \cdot k'} - (-1)^{\alpha \cdot k} \right) 2^{-s} \widehat{S}(\alpha, \beta). \quad (18)$$

Consequently, only Walsh modes satisfying $\alpha \cdot (k \oplus k') = 1$ contribute to the spectral separation. A diagonal local approximation based on the boundary covariance form is

$$d_{\text{loc}}^2(k, k') = \sum_{\substack{(\alpha, \beta) \neq (0, 0): \\ \alpha \cdot (k \oplus k') = 1}} \frac{4 \cdot 2^{-2s} \widehat{S}(\alpha, \beta)^2}{1 - 2^{-2s} \widehat{S}(\alpha, \beta)^2}, \quad (19)$$

provided the denominator is nonzero. This quantity is a local spectral proxy, not the exact Fisher–Rao geodesic distance between the singular distributions.

4 Dual Structure, Boundary Limits, and Belief Geometry

4.1 Regular dual flatness of the interior family

Theorem 4.1 (Regular dual flatness) *The Walsh exponential family (11) is a regular minimal exponential family on $\mathcal{X}$. Its interior carries the standard dually flat information-geometric structure: natural coordinates θ are e -affine, expectation coordinates η are m -affine, the log-partition function ψ is the e -potential, and its Legendre dual is $\varphi(\eta) = -H(p_\eta)$.*

Proof The natural parameter space is the open set $\mathbb{R}^{2^{2s}-1}$. The sufficient statistics are bounded, so the partition function is finite and smooth for all θ . Minimality follows from the orthogonality and linear independence of the nonconstant Walsh characters together with the constant character. The standard theorem for regular minimal exponential families then gives dual flatness and the Legendre correspondence [1, 2]. $\square$

Remark 4.1 (Boundary discipline) Theorem 4.1 applies to full-support distributions only. The exact key distributions P_k lie in the closure. Statements about geodesics, Christoffel symbols, Bregman identities, or curvature at P_k are therefore meaningful only after regularization or after being rephrased as finite boundary covariance statements.

4.2 KL divergence between exact key distributions

Proposition 4.1 (Support separation) *If S is bijective and $k \neq k'$, then*

$$\text{supp}(P_k) \cap \text{supp}(P_{k'}) = \emptyset, \quad D_{\text{KL}}(P_k \| P_{k'}) = +\infty. \quad (20)$$

Proof If (a, b) belongs to both supports, then $S(a \oplus k) = b = S(a \oplus k')$. Since S is injective, $a \oplus k = a \oplus k'$, hence $k = k'$, a contradiction. Thus $P_{k'}$ assigns zero probability to every point in the support of P_k when $k \neq k'$, which implies infinite KL divergence. $\square$

Remark 4.2 (No finite Bregman identity at singular endpoints) The usual equality between KL divergence and the Bregman divergence of the dual potential is an interior statement. For P_k and $P_{k'}$ with disjoint supports, one may study $D_{\text{KL}}(P_k^\varepsilon \| P_{k'}^\varepsilon)$ and then let $\varepsilon \downarrow 0$, but the limiting divergence is infinite. We therefore do not use finite Pythagorean decompositions directly between distinct exact key distributions.

4.3 Mixtures induced by key beliefs

Let $Q = \Delta_{2^s-1}^\circ$ be the interior of the simplex over keys. A belief $q \in Q$ induces a distribution on input-output pairs

$$\bar{P}_q(a, b) = \sum_{k \in \{0,1\}^s} q(k) P_k(a, b) = 2^{-s} \sum_{k \in \{0,1\}^s} q(k) \mathbf{1}[b = S(a \oplus k)]. \quad (21)$$

If q has full support, then $\bar{P}_q$ has full support on $\mathcal{X}$.

Proposition 4.2 (Walsh representation of belief mixtures) *Let*

$$\hat{q}(\alpha) = \sum_{k \in \{0,1\}^s} q(k) (-1)^{\alpha \cdot k}.$$

Then

$$\bar{P}_q(a, b) = 2^{-2s} \sum_{\alpha, \beta \in \{0,1\}^s} \hat{q}(\alpha) \hat{S}(\alpha, \beta) (-1)^{\alpha \cdot a \oplus \beta \cdot b}. \quad (22)$$

In particular, the uniform belief induces the uniform distribution on $\mathcal{X}$.

Proof Apply Walsh inversion to the indicator $\mathbf{1}[b = S(a \oplus k)]$ and sum over k with weights $q(k)$. For uniform q , $\hat{q}(\alpha) = \mathbf{1}[\alpha = 0]$, and $\hat{S}(0, 0) = 2^s$ while $\hat{S}(0, \beta) = 0$ for $\beta \neq 0$ because S is bijective. $\square$

783 4.4 Fisher–Rao geometry of the belief simplex

784

785 **Definition 4.1** (Belief simplex metric) For $q \in Q$ and tangent vectors u, v with $\sum_k u_k =$
 786 $\sum_k v_k = 0$, the Fisher–Rao metric is

$$788 \quad g_q^{(Q)}(u, v) = \sum_{k \in \{0,1\}^s} \frac{u_k v_k}{q(k)}. \quad (23)$$

790

791

792 Under the square-root map $q \mapsto (\sqrt{q(k)})_k$, this metric is the round metric on a
 793 sphere of radius 2. Therefore

794

795

$$796 \quad d_F(q, q') = 2 \arccos \left(\sum_{k \in \{0,1\}^s} \sqrt{q(k)q'(k)} \right). \quad (24)$$

797

798 For distinct point-mass limits, $d_F(\delta_k, \delta_{k'}) = \pi$. This is a statement about the belief
 799 simplex, not about geodesics between the singular input-output distributions P_k and
 800 $P_{k'}$.

801

802

803

804

805

806

807

808

809

810

811

812

813

814

815

816

817

818

819

820

821

822

823

824

825

826

827

828

811 4.5 Pythagorean identities

812 The generalized Pythagorean theorem is used in the following restricted sense. For
 813 full-support distributions in the regular Walsh exponential family, the standard identity

814

815

816

817

818

819

820

821

822

823

824

825

826

827

828

$$818 \quad D_{\text{KL}}(P_A \| P_C) = D_{\text{KL}}(P_A \| P_B) + D_{\text{KL}}(P_B \| P_C) \quad (25)$$

821 holds whenever the e-geodesic from A to B is orthogonal to the m-geodesic from B to
 822 C . In the CPA setting, the relevant full-support object is the posterior belief $q_n \in Q$
 823 as long as the prior has full support and the Gaussian likelihood is strictly positive.
 824 This is why projection statements in Section 5 are made on Q or on regularized joint
 825 distributions, not directly on disjoint-support P_k endpoints.

5 Correlation Power Analysis as Bayesian Projection Dynamics

5.1 Leakage model

We consider the standard non-profiled CPA setting. The adversary knows S , observes plaintexts $a_1, \dots, a_n$ uniformly drawn from $\{0, 1\}^s$, and obtains leakage measurements

$$W_i = f(S(a_i \oplus k^*)) + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2), \quad (26)$$

where k^* is the true key and $f : \{0, 1\}^s \rightarrow \mathbb{R}$ is the leakage model. In the numerical experiments we take $f = \text{HW}$.

5.2 Bayesian update as a mixture projection

Theorem 5.1 (CPA posterior update as m -projection) *Let $q_{i-1} \in \Delta_{2^s-1}^\circ$ be the prior belief over keys before observing (a_i, W_i) . The posterior*

$$q_i(k) = \frac{q_{i-1}(k) p(W_i | a_i, k)}{\sum_{h \in \{0, 1\}^s} q_{i-1}(h) p(W_i | a_i, h)} \quad (27)$$

is the key marginal of the regular conditional distribution obtained from the joint prior-likelihood model. Equivalently, it can be obtained as the limit of m -projections after conditioning on shrinking measurement bins around the observed continuous value.

Proof The joint distribution before conditioning is $q_{i-1}(k)p(a_i)p(W_i | a_i, k)$. Since the Gaussian observation is continuous, the event $W_i = w_i$ has probability zero and conditioning is understood through a regular conditional density, or equivalently as the limit obtained by conditioning on $W_i \in [w_i - \epsilon, w_i + \epsilon]$ and letting $\epsilon \downarrow 0$. The limiting conditional law has key marginal exactly (27). $\square$

The exact S-box distributions P_k lie on the boundary of the simplex, but the belief trajectory $q_0, q_1, \dots, q_n$ stays in the interior of the belief simplex as long as the prior

875 has full support and the Gaussian likelihood is positive. Hence CPA can be studied as
 876
 877 a regular trajectory on $(Q, g_F^{(Q)})$.

878

879

880 5.3 Per-trace signal information

881

882 **Definition 5.1** (Per-trace signal information) For a leakage model f and noise variance σ^2 ,
 883 define

884

$$885 I_S(f, \sigma) = \sigma^{-2} \text{Var}_{a \sim \text{Unif}(\{0,1\}^s)} [f(S(a \oplus k^*))]. \quad (28)$$

886

887 When S is bijective, this quantity is independent of k^* .

888

889

890 **Proposition 5.1** (Hamming-weight specialization) For $f = \text{HW}$ and a bijective S ,

891

$$892 I_S(\text{HW}, \sigma) = \frac{s}{4\sigma^2}. \quad (29)$$

893

894 In particular, for 4-bit bijective S -boxes, $I_S(\text{HW}, \sigma) = 1/\sigma^2$.

895

896

897

898

899 *Proof* Bijectivity of S together with the uniformity of $a \oplus k^*$ implies that $S(a \oplus k^*)$ is uniform
 900 over $\{0, 1\}^s$. The Hamming weight of a uniform s -bit word follows a binomial distribution
 901 Binomial($s, 1/2$), whose variance is $s/4$. □

902

903

904 5.4 Pairwise log-likelihood increments

905

906 For a competing key $k \neq k^*$, the per-trace squared separation is

907

908

909

$$910 \Delta(k, k^*) = \frac{1}{2^s} \sum_{a \in \{0,1\}^s} (f(S(a \oplus k^*)) - f(S(a \oplus k)))^2. \quad (30)$$

911

912

913

914 **Proposition 5.2** (Expected pairwise log-Bayes gain) Under the Gaussian leakage model,

915

$$916 \mathbb{E} \left[\log \frac{p(W | A, k^*)}{p(W | A, k)} \right] = \frac{\Delta(k, k^*)}{2\sigma^2}, \quad (31)$$

917

918 where the expectation is over A (uniform) and W conditioned on A and k^* .

919

920

Proof For a fixed plaintext a , the two likelihoods are Gaussian with the same variance σ^2 and means $\mu_*(a) = f(S(a \oplus k^*))$, $\mu_k(a) = f(S(a \oplus k))$. The Kullback-Leibler divergence between two univariate Gaussians with equal variance is $(\mu_*(a) - \mu_k(a))^2 / (2\sigma^2)$. Averaging over a gives the result. $\square$

The worst-case pairwise gap

$$\Gamma_S(f, \sigma) = \min_{k \neq k^*} \frac{\Delta(k, k^*)}{2\sigma^2} \quad (32)$$

provides a rigorous likelihood-separation measure, distinct from $I_S(f, \sigma)$ in general.

5.5 Model-based sample-complexity estimate

Starting from a uniform prior, the initial key uncertainty equals $s \log 2$ nats. A first-order information estimate then yields

$$n_{\text{est}} \approx \frac{s \log 2}{I_S(f, \sigma)}. \quad (33)$$

For 4-bit Hamming-weight leakage, $n_{\text{est}} \approx 4 \log 2 \sigma^2$. This is a heuristic leakage-rate estimate. A more rigorous bound follows from the pairwise gap $\Gamma_S(f, \sigma)$: the expected log-posterior odds against any fixed wrong key grow at rate at least $\Gamma_S(f, \sigma)$. Thus reaching an expected odds margin τ against the hardest wrong key requires roughly

$$n_{\text{gap}}(\tau) \approx \frac{\tau}{\Gamma_S(f, \sigma)} = \frac{2\sigma^2\tau}{\min_{k \neq k^*} \Delta(k, k^*)}. \quad (34)$$

This model-based estimate links the minimum pairwise gap to a trace budget for expected posterior-odds growth. Two S-boxes with the same scalar I_S can lead to different worst-case posterior-odds growth when their minimum pairwise gaps differ.

967 **6 Active-Eigenspace Curvature: Computation and**
968
969 **Status**
970
971
972 **6.1 What curvature is being computed**
973
974 The regular Walsh exponential family is dually flat with respect to the e- and m-
975 connections, but its Levi-Civita curvature is not automatically zero. At exact key
976 distributions P_k , however, the ordinary Riemannian tensor of the full-support family
977 is not directly defined because P_k lies on the boundary. This section therefore uses the
978 following restricted object.
979
980
981
982
983
984 **Definition 6.1** (Active covariance eigenspace) Let $\mathcal{I}^\partial(k)$ be the boundary covariance matrix
985 in Definition 3.4. Its active eigenspace is the span of the eigenvectors with strictly positive
986 eigenvalues. Curvature computations below are performed after projecting the covariance and
987 third cumulant tensors to this active eigenspace.
988
989
990
991
992 The resulting values are curvature diagnostics of the limiting regularized geometry
993 obtained from the active covariance eigenspace. The finite orbit $\{P_k\}$ itself is discrete
994 and is used here through its induced boundary covariance and cumulant tensors.
995
996
997
998
999 **6.2 Universal rank and eigenvalues of the boundary covariance**
1000
1001 **Theorem 6.1** (Rank and active eigenvalues) *For any bijection $S : \{0, 1\}^s \rightarrow \{0, 1\}^s$, the*
1002 *boundary covariance matrix $\mathcal{I}^\partial(k)$ has rank $2^s - 1$. Its nonzero eigenvalues are all 2^s when the*
1003 *feature coordinates are the unnormalised nonconstant Walsh characters used in Definition 3.3.*
1004
1005
1006
1007
1008
1009 *Proof* Let $N = 2^s$ and index rows by $a \in \{0, 1\}^s$. Let X be the $N \times (2^{2s} - 1)$ feature
1010 matrix with entries $X_{a,(\alpha,\beta)} = (-1)^{\alpha \cdot a \oplus \beta \cdot S(a \oplus k)}$ for $(\alpha, \beta) \neq (0, 0)$. The boundary covariance
1011 is $N^{-1} X^\top H X$, where $H = I_N - N^{-1} \mathbf{1}\mathbf{1}^\top$ centers over the support graph. The nonzero
1012

eigenvalues of $N^{-1}X^\top HX$ equal those of $N^{-1}HXX^\top H$. For $a, a' \in \{0, 1\}^s$,

$$(XX^\top)_{a,a'} = \sum_{(\alpha,\beta) \neq (0,0)} (-1)^{\alpha \cdot (a \oplus a') \oplus \beta \cdot (S(a \oplus k) \oplus S(a' \oplus k))} = N^2 \mathbf{1}[a = a'] - 1,$$

because S is injective. Hence $N^{-1}HXX^\top H = N^{-1}H(N^2I_N - \mathbf{1}\mathbf{1}^\top)H = NH$. The matrix H has rank $N - 1$ and eigenvalue 1 on the centered subspace. Therefore the active eigenvalues are all $N = 2^s$. $\square$

6.3 Third cumulants and curvature protocol

Let

$$C_{ijk}^\partial(k) = \mathbb{E}_{P_k} [(T_i - \eta_i)(T_j - \eta_j)(T_k - \eta_k)] \quad (35)$$

be the boundary third cumulant tensor. Since products of Walsh characters are again Walsh characters, every entry can be computed exactly from the WHT of S and the key-translation rule. If U is an orthonormal basis of the active eigenspace and $\lambda_a = 2^s$ are the active covariance eigenvalues, the projected tensor is

$$\tilde{C}_{abc} = \sum_{i,j,\ell} U_{ia} U_{jb} U_{\ell c} C_{ij\ell}^\partial. \quad (36)$$

The Levi-Civita curvature diagnostic used in the experiments is then computed with the standard dually-flat formula

$$R_{abab} = \frac{1}{4} \sum_c \frac{\tilde{C}_{aac} \tilde{C}_{bbc} - \tilde{C}_{abc}^2}{\lambda_c}, \quad K_{ab} = \frac{R_{abab}}{\lambda_a \lambda_b}, \quad a < b. \quad (37)$$

Both verification scripts implement Equation (37) directly.

1059 6.4 Empirical constant-curvature finding

1060

1061 **Proposition 6.1** (Reproducible 4-bit curvature observation) *For the PRESENT, GIFT, and*
1062 *SKINNY 4-bit S-boxes, the projected active-eigenspace curvature diagnostic (37) gives*
1063

$$1064 \quad K_{ab} = -\frac{1}{4} \quad (38)$$

1066

1067 *for all $\binom{15}{2} = 105$ active coordinate planes, up to numerical precision in the provided scripts.*

1068

1069

1070

1071

1072

1073

1074

1075

1076 *Verification protocol* The computation is reproducible from the repository scripts. The scripts
1077 build the WHT table, construct $\mathcal{I}^\partial$, diagonalize its active 15-dimensional eigenspace, project
1078 the third cumulant tensor, and evaluate (37). The reported value is stable across the three
1079 tested S-boxes and an additional random 4-bit bijection in `verify_curvature.py`.
1080

1081 These reproducible computations suggest the following conjecture: for bijective S-boxes,
1082 the active Fisher subspace associated with the Walsh boundary covariance form has constant
1083 sectional curvature $K = -1/4$, under the curvature convention and projection procedure used
1084 in this work. A symbolic proof, or a characterization of the precise hypotheses under which
1085 the conjecture holds, is left as an open problem. $\square$
1086
1087
1088
1089

1090

1091

1092

1093 6.5 Relation to the belief simplex

1094

1095 The belief simplex $Q = \Delta_{2^s-1}^\circ$ with Fisher–Rao metric is isometric, via the square-root
1096 embedding, to a sphere of radius 2 and hence has constant sectional curvature $+1/4$.
1097
1098

1099 The observed value $-1/4$ in the active boundary covariance diagnostic suggests a sign-
1100 dual phenomenon between the posterior belief geometry and the limiting keyed S-box
1101 covariance geometry, which is recorded here as part of the open geometric structure of
1102 the model.
1103
1104

7 Reproducible Security Diagnostics

The computations in this section connect the closed-form geometry to attack-relevant side-channel diagnostics. First, they verify the exact boundary covariance and active-eigenspace formulas on standard lightweight S-boxes. Second, they separate quantities that are structural invariants from quantities that influence CPA discrimination between key hypotheses. Third, they validate the pairwise-gap interpretation through a reproducible Monte Carlo CPA experiment.

7.1 Target S-boxes

We evaluate the 4-bit bijective S-boxes used in PRESENT, GIFT, and SKINNY [13–15]:

$$S_{\text{PRESENT}} = [C, 5, 6, B, 9, 0, A, D, 3, E, F, 8, 4, 7, 1, 2],$$

$$S_{\text{GIFT}} = [1, A, 4, C, 6, F, 3, 9, 2, D, B, 7, 5, 0, 8, E],$$

$$S_{\text{SKINNY}} = [C, 6, 9, 0, 1, A, 2, B, 3, 8, 5, D, 4, E, 7, F].$$

These examples are useful because they are widely used in lightweight cryptography and because exhaustive enumeration over all inputs, masks, and key hypotheses is possible.

7.2 Walsh spectra and boundary covariance

For each S-box we compute the full WHT table

$$\hat{S}(\alpha, \beta) = \sum_{a \in \{0,1\}^4} (-1)^{\alpha \cdot a \oplus \beta \cdot S(a)}.$$




Fig. 2 Walsh spectra. The plotted Walsh tables show the same amplitude profile for PRESENT, GIFT, and SKINNY, which explains why purely amplitude-based diagnostics coincide for these examples.




Fig. 3 Walsh amplitude distribution. The histogram records the shared counts of absolute Walsh coefficients over nonzero masks for the three tested S-boxes.

The three tested S-boxes have the same absolute Walsh-coefficient distribution over nonzero masks:

$$\#\|\hat{S}\| = 0 : 123, \quad \#\|\hat{S}\| = 4 : 96, \quad \#\|\hat{S}\| = 8 : 36.$$

Therefore diagnostics depending only on the Walsh amplitudes coincide for these three examples. This shared spectral profile provides a useful baseline for reading the later CPA-specific quantities.

The boundary covariance matrix is built from Proposition 3.2. For $s = 4$ it is a 255×255 positive semidefinite matrix. Consistent with Theorem 6.1, all three matrices have rank 15 and active eigenvalues equal to 16 up to numerical precision.

Table 3 Boundary covariance structure for the tested 4-bit bijections.

S-box	Matrix size	Rank	Active eigenvalues
PRESENT	255×255	15	16 repeated 15 times
GIFT	255×255	15	16 repeated 15 times
SKINNY	255×255	15	16 repeated 15 times




Fig. 4 Local key-separation proxy. The matrix visualises the covariance-weighted Walsh displacement between key hypotheses in the local Walsh/Fisher coordinates.

Table 4 Interpretation of the key-separation quantities.

Quantity	Status
$D_{\text{KL}}(P_k \ P_{k'})$ for $k \neq k'$	Infinite, because supports are disjoint
$d_F(\delta_k, \delta_{k'})$ on the belief simplex	Exactly π
$d_{\text{loc}}^2(k, k')$ in (19)	Local Walsh/Fisher covariance proxy

7.3 Local spectral separation

Using (19), we compute local diagonal Walsh/Fisher separation proxies between key hypotheses. This local covariance-weighted coordinate displacement is evaluated in the regularized Walsh/Fisher coordinates introduced earlier. The key-translation rule implies that the value depends only on the key difference $k \oplus k'$.




Fig. 5 Local proxy versus key-difference Hamming weight. The plot aggregates the same covariance-weighted Walsh displacement by the Hamming weight of $k \oplus k'$.

7.4 CPA leakage gaps

For bijective 4-bit S-boxes and uniform inputs, the scalar Hamming-weight variance in Proposition 5.1 is always

$$I_S(\text{HW}, \sigma) = \frac{1}{\sigma^2}.$$

This scalar is a baseline leakage descriptor shared by PRESENT, GIFT, and SKINNY.

Pairwise gaps

$$\Delta(\delta) = \frac{1}{16} \sum_{a \in \{0,1\}^4} (\text{HW}(S(a)) - \text{HW}(S(a \oplus \delta)))^2, \quad \delta \neq 0,$$

are more informative for CPA because the expected log-Bayes gain against a wrong key with difference δ is $\Delta(\delta)/(2\sigma^2)$.

Table 5 shows how the scalar variance descriptor and the pairwise structure play complementary roles: all three S-boxes have the same average pairwise gap, while the worst and best key-difference cases vary. These pairwise gaps are the quantities that enter expected posterior-odds growth in the Gaussian leakage model.

Table 5 Pairwise Hamming-weight leakage gaps over nonzero key differences.

S-box	$\min_{\delta \neq 0} \Delta(\delta)$	$\max_{\delta \neq 0} \Delta(\delta)$	$\frac{1}{15} \sum_{\delta \neq 0} \Delta(\delta)$
PRESENT	1	7/2	32/15
GIFT	5/4	7/2	32/15
SKINNY	1	3	32/15

Table 6 Operational CPA interpretation of the pairwise gaps. The listed deltas are nonzero key differences $\delta = k \oplus k^*$ in hexadecimal.

S-box	Hardest min Δ	gaps Easiest max Δ	gaps Gap multiset over $\delta \neq 0$
PRESENT	1 at $\{A, E\}$	7/2 at $\{6\}$	$1^2; (5/4)^1; (3/2)^2; (7/4)^3; (5/2)^2;$ $(11/4)^1; 3^2; (13/4)^1; (7/2)^1$
GIFT	5/4 $\{C, E, F\}$	at 7/2 at $\{7\}$	$(5/4)^3; (3/2)^2; 2^2; (9/4)^3; (5/2)^2; 3^2;$ $(7/2)^1$
SKINNY	1 at $\{5\}$	3 at $\{B\}$	$1^1; (5/4)^1; (3/2)^1; (7/4)^1; 2^3; (9/4)^2;$ $(5/2)^3; (11/4)^2; 3^1$

Table 6 gives the CPA-oriented reading of the pairwise theory. PRESENT and SKINNY have a smaller hardest-key gap than GIFT under this leakage model, so the model-based estimate (34) assigns a larger trace budget to the hardest posterior-odds margin for those S-boxes. The comparison is local to the specified S-box computation and leakage model, which is the level at which CPA key hypotheses are evaluated.

7.5 Monte Carlo CPA validation

To connect the gap analysis to an attack observable, we simulate the Bayesian CPA model of Section 5. For each S-box, we use true key $k^* = 0$, uniformly sampled plaintext nibbles, leakage $W = \text{HW}(S(a \oplus k^*)) + \varepsilon$, Gaussian noise standard deviation $\sigma = 2$, and 3000 Monte Carlo trials with seed 20260622. For every trace prefix, all 16 key log-likelihoods are accumulated and the true-key rank is recorded. The script `simulate_cpa.py` regenerates both Fig. 6 and the CSV summary `cpa_simulation_results.csv`.




Fig. 6 Simulated CPA under Hamming-weight Gaussian leakage. The left panel shows empirical success probability and the right panel shows mean true-key rank over 3000 trials. The curves translate the pairwise-gap analysis into trace-dependent attack observables.




Fig. 7 Per-trace leakage baseline. The scalar Hamming-weight variance is the same for all bijective 4-bit S-boxes under uniform inputs, so CPA discrimination is completed by pairwise leakage gaps that determine key-hypothesis discrimination.

The simulation confirms the qualitative role of the pairwise gaps. At 64 traces, the empirical success probabilities are approximately 0.923 for PRESENT, 0.950 for GIFT, and 0.947 for SKINNY; at 128 traces, all three success probabilities are between 0.993 and 0.996. Mean true-key ranks decrease toward one for all three S-boxes. The finite-trace curves summarize the combined effect of all wrong-key competitors, while the minimum gap in Table 5 controls the hardest pairwise posterior-odds growth.

Table 7 Projected active-eigenspace curvature diagnostic for the tested 4-bit S-boxes.

S-box	Active dimension	Planes	Min K	Max K
PRESENT	15	105	−0.25	−0.25
GIFT	15	105	−0.25	−0.25
SKINNY	15	105	−0.25	−0.25

7.6 Active-eigenspace curvature

The curvature diagnostic of Section 6 is evaluated in the 15-dimensional active eigenspace of the boundary covariance matrix. The result is constant for all coordinate planes in the tested cases.

7.7 Reproducibility checklist

The computational claims in this section are reproducible from the additional files. The scripts `gen_figures.py`, `simulate_cpa.py`, `compute_curvatures.py`, and `verify_curvature.py` regenerate the figures, CPA simulation results, active covariance diagnostics, and curvature checks. The rank/eigenvalue statement is proved independently in Theorem 6.1; the scripts provide finite checks for the reported 4-bit examples.

8 Discussion

8.1 What the framework establishes

The framework establishes a controlled information-geometric model for the local S-box distributions that appear in side-channel key recovery. Its first role is structural: exact key-induced distributions are singular boundary points, while ordinary Fisher–Rao geometry belongs to full-support interiors or to the adversary’s belief simplex. This distinction prevents assigning non-degenerate Fisher metrics, finite KL divergences, or ordinary geodesics to disjoint-support key endpoints.

1427 The second role is algebraic. Walsh characters give explicit coordinates for the
1428 boundary orbit, and the key-translation rule determines how expectation coordinates
1429 move when the key hypothesis changes. This rule is elementary, but it is useful because
1430 it gives closed-form covariance and displacement expressions.
1431

1432 The third role is security-oriented. CPA is a Bayesian trajectory on the belief
1433 simplex. Under the Gaussian leakage model, pairwise Hamming-weight gaps determine
1434 expected log-Bayes gains against wrong keys. The scalar leakage variance is therefore
1435 a baseline descriptor, whereas pairwise gaps describe the discrimination structure
1436 relevant to key recovery.
1437

1438 This gives a concrete review criterion for the proposed quantities. A diagnostic
1439 is useful for side-channel security when it changes the interpretation of posterior-
1440 odds growth, hard key differences, or the validity of a geometric approximation. The
1441 numerical section combines the shared Walsh-amplitude invariants of PRESENT,
1442 GIFT, and SKINNY with the pairwise gap distributions and simulated CPA curves that
1443 determine which wrong keys are hardest to eliminate under the stated leakage model.
1444

1445

1446 8.2 Position within S-box and side-channel analysis

1447

1448 The framework uses established 4-bit S-box classifications as background and focuses
1449 on the local distributions used in side-channel key recovery. PRESENT, GIFT, and
1450 SKINNY serve as reproducible lightweight examples with shared Walsh-amplitude
1451 profiles and different CPA gap distributions. The standard Gaussian Hamming-weight
1452 CPA model provides the inference setting in which the boundary S-box law, posterior
1453 belief law, and pairwise discrimination gaps are connected.
1454

1455

1456 8.3 Cryptographic interpretation

1457

1458 For isolated bijective 4-bit S-boxes under uniform input and Hamming-weight leakage,
1459 $I_S(\text{HW}, \sigma) = 1/\sigma^2$ for every bijection. This equality is best used as a baseline descriptor.
1460

Practical side-channel behavior is shaped by pairwise leakage gaps, implementation	1473
leakage, noise, masking, alignment, higher-order effects, and the full cipher context.	1474
	1475
In this sense, identical Walsh-amplitude histograms are a useful baseline. They show	1476
which quantities are shared by the three tested S-boxes, while the CPA gap distribution	1477
and Monte Carlo curves expose model-dependent hard and easy key differences. This	1478
is precisely the separation required before using information-geometric language in a	1479
cybersecurity setting.	1480
	1481
The boundary covariance and local Walsh/Fisher separation proxy complement	1482
LAT, DDT, and empirical leakage evaluation. They provide coordinate diagnostics	1483
for the singular key orbit and inputs to posterior belief dynamics alongside standard	1484
cryptanalytic criteria.	1485
	1486
	1487
	1488
	1489
	1490
	1491
	1492
	1493
8.4 Status of the curvature finding	1494
	1495
The active-eigenspace value $K = -1/4$ is reported as a reproducible diagnostic under	1496
the projection, normalization, and sign conventions specified in Section 6. The rank and	1497
active eigenvalue statement is proved for every bijective S-box. A complete symbolic	1498
characterization of the curvature diagnostic, including the exact hypotheses under	1499
which it is constant, remains open.	1500
	1501
	1502
	1503
	1504
	1505
	1506
8.5 Open problems	1507
	1508
Several extensions are natural. First, the $\varepsilon \downarrow 0$ limit of regularized Fisher–Rao geometry	1509
could be studied using tools from singular information geometry. Second, the active-	1510
eigenspace curvature phenomenon could be attacked symbolically. Third, the leakage	1511
model could be extended to masked, higher-order, non-Gaussian, or profiled settings.	1512
	1513
Fourth, full-round ciphers can be handled by adding propagation through linear layers	1514
and key schedules.	1515
	1516
	1517
	1518

1519 9 Conclusion

1520

1521 We presented a boundary information-geometric framework for the local S-box dis-
1522 tributions that arise in side-channel key recovery. The main modelling point is the
1523 separation between three spaces: the regular full-support Walsh exponential family,
1524 the singular key-induced boundary orbit, and the adversary’s belief simplex. This
1525 separation makes Fisher, KL, and projection statements mathematically well scoped.

1526 For bijective S-boxes, Walsh coordinates give closed-form boundary covariance
1527 matrices. The active covariance subspace has dimension $2^s - 1$ and all active eigenvalues
1528 equal 2^s in the unnormalised Walsh basis. Thus this part of the geometry is a general
1529 structural consequence of bijectivity.

1530 For CPA, Bayesian updates occur on the belief simplex. Under Hamming-weight
1531 Gaussian leakage, the scalar variance information is a baseline fixed by bijectivity,
1532 while pairwise leakage gaps determine expected posterior-odds growth against wrong
1533 keys. Computations on PRESENT, GIFT, and SKINNY illustrate this distinction and
1534 reproduce the active curvature diagnostic $K = -1/4$ in the tested 4-bit cases.

1544 The framework contributes a controlled geometric language for connecting singular
1545 S-box distributions, boundary covariance, and side-channel Bayesian inference. The
1546 security value of the framework is to make explicit which geometric quantities are
1547 structural invariants and which quantities affect local key-hypothesis discrimination
1548 under a specified leakage model.

1549

1550

1551

1552

1553 Declarations

1554

1555

1556 Availability of data and material

1557

1558 The datasets supporting the conclusions of this article are included within the arti-
1559 cle and its additional files. Additional file 1, `supplementary_reproducibility.pdf`,
1560 records the computational protocol, software environment, arithmetic conventions, and
1561

reproducibility coverage. Additional file 2, `supplementary_material.zip`, contains the scripts, dependency list, CPA simulation CSV, and reproduction README used to regenerate the numerical tables and figures.

Software information: project name, Boundary Information Geometry for S-box Side-Channel Leakage; project home page, <https://anonymous.4open.science/r/IG-based-Sbox-SCL-DE49/>; operating system, platform independent; programming language, Python 3; other requirements, `numpy`, `scipy`, and `matplotlib`; license, not specified; restrictions to use by non-academics, none beyond the repository license status.

Funding

Not applicable.

Acknowledgements

Not applicable.

References

- [1] Amari, S.-i., Nagaoka, H.: Methods of Information Geometry. Translations of Mathematical Monographs, vol. 191. American Mathematical Society / Oxford University Press, Providence, RI / Oxford (2000). Translated by Daishi Harada
- [2] Ay, N., Jost, J., Lê, H.V., Schwachhöfer, L.: Information Geometry. Ergebnisse der Mathematik und ihrer Grenzgebiete. Springer, Cham (2017)
- [3] Chentsov, N.N.: Statistical Decision Rules and Optimal Inference. Translations of Mathematical Monographs, vol. 53. American Mathematical Society, Providence, RI (1982)

- 1611 [4] Ikeda, S., Tanaka, T., Amari, S.-i.: Information geometry of turbo and low-density
1612 parity-check codes. *IEEE Transactions on Information Theory* **50**(6), 1097–1114
1613 (2004) <https://doi.org/10.1109/TIT.2004.828072>
1614
1615
1616
1617 [5] Nyberg, K.: Differentially uniform mappings for cryptography. In: *Advances*
1618 *in Cryptology – EUROCRYPT 1993*. *Lecture Notes in Computer Science*, vol.
1619 765, pp. 55–64. Springer, Berlin, Heidelberg (1994). [https://doi.org/10.1007/](https://doi.org/10.1007/3-540-48285-7_6)
1620 [3-540-48285-7_6](https://doi.org/10.1007/3-540-48285-7_6)
1621
1622
1623
1624 [6] Leander, G., Poschmann, A.: On the classification of 4 bit S-boxes. In: *Arith-*
1625 *metic of Finite Fields – WAIFI 2007*. *Lecture Notes in Computer Science*, vol.
1626 4547, pp. 159–176. Springer, Berlin, Heidelberg (2007). [https://doi.org/10.1007/](https://doi.org/10.1007/978-3-540-73074-3_13)
1627 [978-3-540-73074-3_13](https://doi.org/10.1007/978-3-540-73074-3_13)
1628
1629
1630
1631
1632 [7] Beyne, T.: A geometric approach to linear cryptanalysis. In: *Advances in Cryptol-*
1633 *ogy – ASIACRYPT 2021*. *Lecture Notes in Computer Science*, vol. 13090, pp.
1634 36–66. Springer, Cham (2021). https://doi.org/10.1007/978-3-030-92068-5_2
1635
1636
1637
1638 [8] Beyne, T., Rijmen, V.: Differential cryptanalysis in the fixed-key model. In:
1639 *Advances in Cryptology – CRYPTO 2022*. *Lecture Notes in Computer Sci-*
1640 *ence*, vol. 13509, pp. 687–716. Springer, Cham (2022). [https://doi.org/10.1007/](https://doi.org/10.1007/978-3-031-15982-4_23)
1641 [978-3-031-15982-4_23](https://doi.org/10.1007/978-3-031-15982-4_23)
1642
1643
1644
1645 [9] Kocher, P., Jaffe, J., Jun, B.: Differential power analysis. In: *Advances in Cryptol-*
1646 *ogy – CRYPTO 1999*. *Lecture Notes in Computer Science*, vol. 1666, pp. 388–397.
1647 Springer, Berlin, Heidelberg (1999). https://doi.org/10.1007/3-540-48405-1_25
1648
1649
1650
1651 [10] Brier, E., Clavier, C., Olivier, F.: Correlation power analysis with a leakage model.
1652 In: *Cryptographic Hardware and Embedded Systems – CHES 2004*. *Lecture Notes*
1653 *in Computer Science*, vol. 3156, pp. 16–29. Springer, Berlin, Heidelberg (2004).
1654
1655
1656

- https://doi.org/10.1007/978-3-540-28632-5_2 1657
1658
- [11] Standaert, F.-X., Malkin, T.G., Yung, M.: A unified framework for the analysis 1659
of side-channel key recovery attacks. In: Advances in Cryptology – EUROCRYPT 1660
2009. Lecture Notes in Computer Science, vol. 5479, pp. 443–461. Springer, Berlin, 1661
Heidelberg (2009). https://doi.org/10.1007/978-3-642-01001-9_26 1662
1663
1664
1665
1666
- [12] Veyrat-Charvillon, N., Standaert, F.-X.: Mutual information analysis: How, when 1667
and why? In: Cryptographic Hardware and Embedded Systems – CHES 2009. 1668
Lecture Notes in Computer Science, vol. 5747, pp. 429–443. Springer, Berlin, 1669
Heidelberg (2009). https://doi.org/10.1007/978-3-642-04138-9_30 1670
1671
1672
1673
- [13] Bogdanov, A., Knudsen, L.R., Leander, G., Paar, C., Poschmann, A., Robshaw, 1674
M.J.B., Seurin, Y., Vikkelsoe, C.: PRESENT: An ultra-lightweight block cipher. 1675
In: Cryptographic Hardware and Embedded Systems – CHES 2007. Lecture Notes 1676
in Computer Science, vol. 4727, pp. 450–466. Springer, Berlin, Heidelberg (2007). 1677
https://doi.org/10.1007/978-3-540-74735-2_31 1678
1679
1680
1681
1682
- [14] Banik, S., Pandey, S.K., Peyrin, T., Sasaki, Y., Sim, S.M., Todo, Y.: GIFT: A small 1683
present – towards reaching the limit of lightweight encryption. In: Cryptographic 1684
Hardware and Embedded Systems – CHES 2017. Lecture Notes in Computer 1685
Science, vol. 10529, pp. 321–345. Springer, Cham (2017). https://doi.org/10.1007/978-3-319-66787-4_16 1686
1687
1688
1689
1690
1691
- [15] Beierle, C., Jean, J., Kölbl, S., Leander, G., Moradi, A., Peyrin, T., Sasaki, Y., 1692
Sasdrich, P., Sim, S.M.: The SKINNY family of block ciphers and its low-latency 1693
variant MANTIS. In: Advances in Cryptology – CRYPTO 2016. Lecture Notes in 1694
Computer Science, vol. 9815, pp. 123–153. Springer, Berlin, Heidelberg (2016). 1695
https://doi.org/10.1007/978-3-662-53008-5_5 1696
1697
1698
1699
1700
1701
1702